

13. Dijkstra's

Kristiāns Vinters

Fall 2023

Introduction

I solved the assignment in Go. I used Go because I want to become more familiar with it. Source code and benchmark data is available on GitHub¹.

Implementation

For the implementation, I reused the `trains` package from the graphs assignment and the `pqueue` package from the heap assignment. I did not need to make any significant modifications to either package for this assignment. The only modification I made to the `trains` package is that I added the `idx` field to the `City` struct.

```
type City struct {
    Name      string
    idx       int // Sequential city index set when loading data
    Neighbors []Connection
}
```

The Dijkstra pathfinding algorithm was implemented as the `DijkstraFind()` method on the `NetworkGraph` struct. The implementation of the algorithm is explained with code comments.

```
func (n *NetworkGraph) DijkstraFind(
    from,
    to string
) (int, []string, error) {
    // Exit early if the from and to cities are the same
    if from == to {
        return 0, []string{from}, nil
    }
}
```

¹<https://github.com/Phanty133/id1021/tree/master/13-dijkstra>

```

// Get the City objects from the graph by name
fromCity := n.LookupCity(from)
toCity := n.LookupCity(to)

if fromCity == nil {
    return 0, nil, fmt.Errorf("city %s not found", from)
}

if toCity == nil {
    return 0, nil, 0, fmt.Errorf("city %s not found", to)
}

// Initialize both data structures
pq := pqueue.NewArrHeap[*PathNode](n.NumCities)
cityNodes := make([]*PathNode, n.NumCities)

// Add the starting city to the priority queue, with distance=0
pq.Add(&PathNode{fromCity, nil, 0, 0}, 0)

// Iterate until all cities have been visited
for !pq.Empty() {
    cur, _ := pq.Dequeue()

    // If the current city is the destination, we're done
    if cur.city == toCity {
        // Assemble the path we found and return
        path := make([]string, 0)
        city := cur.city

        for city != nil {
            path = append(path, city.Name)
            city = cityNodes[city.idx].prev
        }

        return cur.dist, path, nil
    }

    // Make sure the current city's Node is counted as visited
    cityNodes[cur.city.idx] = cur

    for _, n := range cur.city.Neighbors {
        node := cityNodes[n.City.idx]

        // If the neighbor city hasn't been visited yet,

```

```

        // add it to the queue
        if node == nil {
            node = &PathNode{n.City, cur.city, n.Dist + cur.dist, 0}
            cityNodes[n.City.idx] = node
            pq.Add(node, node.dist)
        }

        // If the neighbor city has been visited,
        // but the current path is shorter,
        // update the neighbor city's Node
        // and update the priority queue
        if node.dist > n.Dist+cur.dist {
            node.dist = n.Dist + cur.dist
            node.prev = cur.city
            pq.Update(node, node.dist)
        }
    }
}

return 0, nil, errors.New("no path found")
}

```

Benchmarks

I benchmarked Dijkstra’s algorithm first by searching for the same routes as in the graphs assignment on the `trains.csv` dataset, and then by searching for routes from Berlin to all other cities in the `europe.csv` dataset. All searches were performed 100 times, and the median time was used for the results.

Comparison with previous best

To compare with the previous best implementation from the graphs assignment, I ran the same routes with Dijkstra’s algorithm. The results are shown in Figure 1.

From	To	Dist, min	Previous best, μs	Dijkstra's, μs
Göteborg	Stockholm	211	58.4	2.86
Göteborg	Umeå	705	28150	5.24
Göteborg	Sundsvall	515	5100	4.89
Malmö	Göteborg	153	61.4	1.89
Malmö	Stockholm	273	31.0	2.44
Malmö	Kiruna	1162	272000	5.80
Stockholm	Sundsvall	327	2770	2.86
Stockholm	Umeå	517	10180	4.26
Sundsvall	Umeå	190	812500	1.89
Umeå	Göteborg	705	108	4.26

Figure 1: Pathfinding results

The times show that Dijkstra's algorithm is significantly faster for all routes, as well as being more consistent. A contributing factor for the reason that the previous best is so inconsistent in its times is that it has to do a linear search over the path visited for every new city it visits, therefore longer routes like the one from Malmö to Kiruna take significantly longer than shorter routes like the one from Malmö to Göteborg.

Europe graph

I benchmarked Dijkstra's algorithm on the `europe.csv` dataset by searching for routes from Berlin to all other cities. I then plotted the results vs route time, route node count, and number of cities visited during the search. The results are shown in Figures 2, 3, and 4 respectively.




Figure 2: Median lookup time vs route time

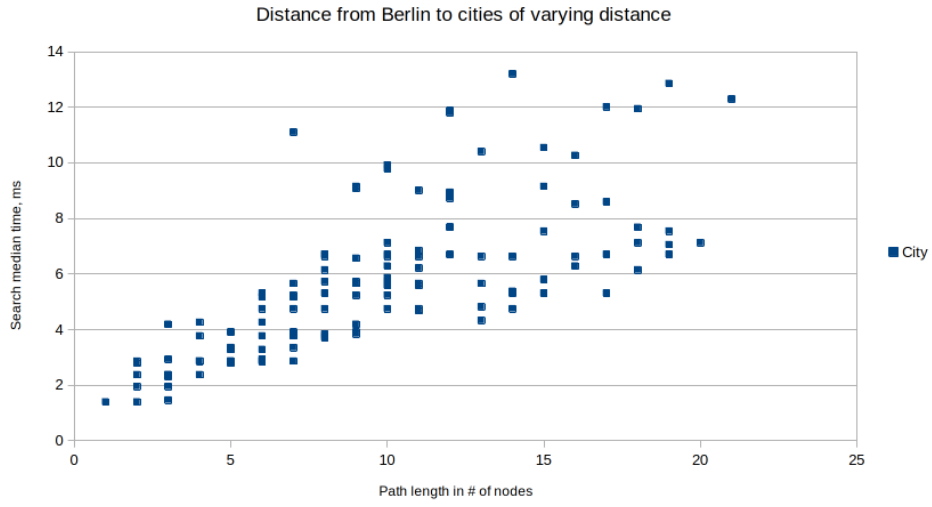


Figure 3: Median lookup time vs route distance in # of nodes




Figure 4: Median lookup time vs # of nodes visited

Figures 2 and 3 show an evident upwards trend in lookup time as route time and route node count increase, however both have a lot of spread in the distribution. Figure 4 shows a clearer, almost linear trend than the other two figures with a tighter spread.